

SDS323 Statistical Learning And Inference
Final Project
Nitrogen Dioxide: Air Pollution And Road Traffic At Alnabru In
Oslo, Norway
Lazar Ilic
EID li948

This project was produced entirely by Lazar Ilic. In a 1 session
600 minute round setting on 2022-05-08.

Introduction

The data are a subsample of 500 observations from a data set that originate in a study where air pollution at a road is related to traffic volume and meteorological variables, collected by the Norwegian Public Roads Administration. The response variable [column 1] consist of hourly values of the logarithm of the concentration of NO2 [particles], measured at Alnabru in Oslo, Norway, between October 2001 and August 2003. The predictor variables [columns 2 to 8] are the logarithm of the number of cars per hour, temperature 2 meter above ground [degree C], wind speed [meters/second], the temperature difference between 25 and 2 meters above ground [degree C], wind direction [degrees between 0 and 360], hour of day and day number from Monday October 1. 2001. Submitted by Magne Aldrin [magne.aldrin@nr.no]. [28/Jul/04] [19kbytes]

I suppose that I might try to simply predict NO2 from the measurements of the other variables. Or perhaps try to predict the number of cars per hour from the other variables. It is not totally clear to me how costly each of those 2 measurements are or what the sort of latency involved in such measurements is. How useful these metrics are for other downstream ones. In any case a simple prediction task, running validation to measure model performance, and try to produce a good simple model. Perhaps run a subset selection algorithm or call the nls and glm functions in .R too maybe to demonstrate just how much estimated performance can improve in a variety of settings and discuss the computational aspects of these optimisations and the perhaps gradient descents involved in finding the optimal coefficients to minimise error.

Methods

The 1st natural subtask is to simply ideate a little bit. I had previously done a sort of task kind of like this one. With only 500 observations we aim to produce a relatively simple model, maybe 5 or 6 derivative predictor variables, all with strong validity and t-scores. Here, we consider exponentiating certain variables and re transforming them in other more useful ways. As in last time, I considered the day of the week. This is because I expect a different distribution of car types for weekends and weekdays and the number of cars per hour is a very very low resolution metric. Also one might think that there are non trivial sources of air pollution from roads in certain directions, or other potential sources and so bucketing and modifying based upon wind direction may be useful. A natural physics interpretation of air movement and rate of air transfer and throughput off of the road level has to do with the measured temperature values and the Delta which can signify a difference in air pressure or simply reflect facets of the measurement location which I cannot discern simply from afar in this dataset.

A 2nd subtask may be to consult some of the extant literature on this and similar related topics. Here of course we have been engaged in studying the most fundamental basics of statistics from the book An Introduction to Statistical Learning: With Applications in R. But there are papers on Air Pollution and one can query Google, ArXiv, Google Scholar, etc. with "NO2 prediction", "NO2 statistics", "air pollution prediction", etc.

Artificial Intelligence Accuracy Assessment In NO2 Concentration Forecasting Of Metropolises Air

NO2 Prediction Using Machine Learning Analyses In Google EE
Prediction Of The NO2 Concentration Data In An Urban Area Using Multiple Regression And Neuronal Networks

Estimating Lockdown-Induced European NO2 Changes: Using Satellite And Surface Observations And Air Quality Models

A Novel Method For Regional NO2 Concentration Prediction Using Discrete Wavelet Transform And An Long Short Term Memory Network
Spatialisation And Prediction Of Seasonal NO2 Pollution Due To Climate Change In The Korean Capital Area Through Land Use Regression Modeling

Error Prediction Of Air Quality At Monitouring Stations Using Random Forest In A Total Error Framework

Nitrogen Dioxide Levels Estimated From Land Use Regression Models Several Years Apart And Association With Mortality In A Large Cohort Study

A 3rd subtask is to execute a semi comprehensive Exploratory Data Analysis. This comes at the beginning of pretty much every single strong Kaggle Notebook, Project, and Competition submission. So one can produce plots and think about models say with simple subsets of the predictors first like 1 predictor, 2 predictors, plausible interaction terms and study these in isolation and view plots of residuals as in our earlier weeks to examine potential functions for relations. These are reasonable methods.

The 4th subtask is to actually muck around, get our hands dirty, and produce some models for predicting NO2 from the other predictors. Run Cross Validation and Root Mean Squared Error metric as in class to test out our various models and pick one which seems best from our perspective. Perhaps not quite the most performant will be our centaur bet on the actually most useful, valid, and human legible model. We might be prone to overly hacking a little bit and if we try out too many models, the most complex one which produces a strong performance may not really be the best, it may be due to a sort of statistical overfitting or just getting lucky on some pseudorandom complex bullpoopoo. The code will reveal these techniques but it suffices to say that some of the tricks we pulled in class as well as perhaps nls, glm, and other tricks from web logs on these topics come in handy.

The 5th subtask if I get around to it is to predict the number of cars per hour from the other variables. But otherwise it is rather similar to the 4th subtask.

Conclusion And Future Work

Summary Of Results

We produced a simple linear regression model with an R^2 of for the provided logarithm of NO2 concentration variable. This corresponds with a Mean Squared Error of on an average variable of 0.5323 on 5 inputs. So that seems fairly good.

Call:

```
lm(formula = logNO2 ~ logcars + temp + winds + hour + daysplit,  
    data = z)
```

Residuals:

	Min	1Q	Median	3Q	Max
	-2.09434	-0.30500	-0.00121	0.36370	1.85733

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)	
(Intercept)	1.368052	0.172938	7.911	1.69e-14	***
logcars	0.444086	0.027539	16.126	< 2e-16	***
temp	-0.026068	0.003776	-6.903	1.57e-11	***
winds	-0.147835	0.013605	-10.866	< 2e-16	***
hour	-0.013929	0.004305	-3.236	0.00129	**
daysplit	-0.292934	0.048788	-6.004	3.73e-09	***

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 0.5323 on 494 degrees of freedom
Multiple R-squared: 0.5022, Adjusted R-squared: 0.4971
F-statistic: 99.66 on 5 and 494 DF, p-value: < 2.2e-16

Along the way, a number of interesting sub results and observations were made.

Potential Future Work

More data. More information about measurements, instruments, and their errors. More precise geographic data. Ideas about how to modify models given say upcoming future construction information and projects in the region. When someone pings a city for a permit, how far in advance can we use this to predict something useful? How complex are useful models for mapping out future growth of cities trajectories and downstream air pollution dynamics? Will the city respond with a strong tax to disincentivise and thereby lower air pollution? Will the people respond by trying harder to move to suburbs or higher floor apartments? Study the histories of big cities and maybe learn something, London, Los Angeles, Tokyo, New York City, Beijing, Hong Kong, etc. When the pandemic hit how fast were trading firms to predict a variety of trajectories of a variety of different things and how those downstream effects would impact a variety of sectors especially with respect to mid and low frequency trading? Did Nassim Nicholas Taleb make a ton of money shorting a ton back before the pandemic really erupted but he called and predicted that it would in fact blow up to massive epic proportion?

```
# SDS323 Statistical Learning And Inference
# Final Project
# Lazar Ilic
# EID li948
# Sources Referenced: 2 Github R Tutorials, Blogs, and R Documentation

library(tidyverse)
library(tm)
library(ISLR)
#library(ISLR2)
library(e1071)
library(pls)
library(MASS)
library(class)
library(glmnet)
library(leaps)
library(caTools)
library(splines)
library(boot)
library(mgcv)
library(tree)
library(randomForest)
library(gbm)

z=read.csv("../input/inferencedataset/Final.csv")

# "logNO2", "logcars", "temp", "winds", "tempd", "windd", "hour", "day"

View(z)
summary(z)

plot(z$logcars, z$logNO2, xlab="logcars", ylab="logNO2", main="logNO2 Versus logcars")
plot(z$temp, z$logNO2, xlab="temp", ylab="logNO2", main="logNO2 Versus temp")
plot(z$winds, z$logNO2, xlab="winds", ylab="logNO2", main="logNO2 Versus winds")
plot(z$tempd, z$logNO2, xlab="tempd", ylab="logNO2", main="logNO2 Versus tempd")
plot(z$windd, z$logNO2, xlab="windd", ylab="logNO2", main="logNO2 Versus windd")
plot(z$hour, z$logNO2, xlab="hour", ylab="logNO2", main="logNO2 Versus hour")
plot((z$day)%%7, z$logNO2, xlab="day", ylab="logNO2", main="logNO2 Versus day")
plot((z$day)%%365, z$logNO2, xlab="day365", ylab="logNO2", main="logNO2 Versus day365")

# It immediately strikes the viewer that logcars looks to be a relatively useful predictor

# "logNO2", "logcars", "temp", "winds", "tempd", "windd", "hour", "day"

z$wind1=ifelse(140<=z$windd & z$windd<=360,1,0)
#View(z)
z$daytime=ifelse(8<=z$hour & z$hour<=18,1,0)
z$day7=(z$day)%%7
z$day365=(z$day)%%365
z$daysplit=ifelse(z$day7<1.5 | z$day7>5.5,1,0)
```

```
View(z)
```

```
# Do a little bit of the old un hacked eyeballed binary Boolean one hot dummy variable cor
```

```
# "logNO2", "logcars", "temp", "winds", "tempd", "windd", "hour", "day", "wind1", "daytime", "day
```

```
Model.logcars=lm(logNO2 ~ logcars, z)
```

```
summary(Model.logcars)
```

```
plot(Model.logcars$fitted.values, Model.logcars$residuals, main="Residuals For Model.logcars")
```

```
hist(Model.logcars$residuals)
```

```
Model.temp=lm(logNO2 ~ temp, z)
```

```
summary(Model.temp)
```

```
plot(Model.temp$fitted.values, Model.temp$residuals, main="Residuals For Model.temp", ylab="F
```

```
hist(Model.temp$residuals)
```

```
Model.winds=lm(logNO2 ~ winds, z)
```

```
summary(Model.winds)
```

```
plot(Model.winds$fitted.values, Model.winds$residuals, main="Residuals For Model.winds", ylab="F
```

```
hist(Model.winds$residuals)
```

```
Model.tempd=lm(logNO2 ~ tempd, z)
```

```
summary(Model.tempd)
```

```
plot(Model.tempd$fitted.values, Model.tempd$residuals, main="Residuals For Model.tempd", ylab="F
```

```
hist(Model.tempd$residuals)
```

```
Model.windd=lm(logNO2 ~ windd, z)
```

```
summary(Model.windd)
```

```
plot(Model.windd$fitted.values, Model.windd$residuals, main="Residuals For Model.windd", ylab="F
```

```
hist(Model.windd$residuals)
```

```
# "logNO2", "logcars", "temp", "winds", "tempd", "windd", "hour", "day", "wind1", "daytime", "day
```

```
Model.hour=lm(logNO2 ~ hour, z)
```

```
summary(Model.hour)
```

```
plot(Model.hour$fitted.values, Model.hour$residuals, main="Residuals For Model.hour", ylab="F
```

```
hist(Model.hour$residuals)
```

```
Model.wind1=lm(logNO2 ~ wind1, z)
```

```
summary(Model.wind1)
```

```
plot(Model.wind1$fitted.values, Model.wind1$residuals, main="Residuals For Model.wind1", ylab="F
```

```
hist(Model.wind1$residuals)
```

```
Model.daytime=lm(logNO2 ~ daytime, z)
```

```
summary(Model.daytime)
```

```
plot(Model.daytime$fitted.values, Model.daytime$residuals, main="Residuals For Model.daytime", ylab="F
```

```
hist(Model.daytime$residuals)
```

```
Model.day7=lm(logNO2 ~ day7, z)
```

```
summary(Model.day7)
```

```
plot(Model.day7$fitted.values, Model.day7$residuals, main="Residuals For Model.day7", ylab="F
```

```
hist(Model.day7$residuals)

Model.day365=lm(logNO2 ~ day365,z)
summary(Model.day365)
plot(Model.day365$fitted.values,Model.day365$residuals,main="Residuals For Model.day365",)
hist(Model.day365$residuals)

Model.daysplit=lm(logNO2 ~ daysplit,z)
summary(Model.daysplit)
plot(Model.daysplit$fitted.values,Model.daysplit$residuals,main="Residuals For Model.daysplit",)
hist(Model.daysplit$residuals)

# "logNO2", "logcars", "temp", "winds", "tempd", "windd", "hour", "day", "wind1", "daytime", "day7",

# logcars seems useful. temp seems useful. winds seems useful. hour seems useful. daytime

Model1=lm(logNO2 ~ logcars + temp + winds + hour + daytime + daysplit,z)
summary(Model1)
plot(Model1$fitted.values,Model1$residuals,main="Residuals For Model1",ylab="Residual",xlab="logcars",)
hist(Model1$residuals)

# OK now a huge loss of significance on daytime.

Model2=lm(logNO2 ~ logcars + temp + winds + hour + daysplit,z)
summary(Model2)
plot(Model2$fitted.values,Model2$residuals,main="Residuals For Model2",ylab="Residual",xlab="logcars",)
hist(Model2$residuals)

# Could try more sophisticated things like glm, nls. Or forming more complex models and run
```



— Attaching packages — tidyverse 1.3.1 —

✓ ggplot2 3.3.5	✓ purrr 0.3.4
✓ tibble 3.1.5	✓ dplyr 1.0.7
✓ tidyr 1.1.4	✓ stringr 1.4.0
✓ readr 2.0.2	✓ forcats 0.5.1

— Conflicts — tidyverse_conflicts() —

✗ dplyr::filter() masks stats::filter()
✗ dplyr::lag() masks stats::lag()

Loading required package: NLP

Attaching package: 'NLP'

The following object is masked from 'package:ggplot2':

annotate

The following object is masked from 'package:httr':

content

Attaching package: 'pls'

The following object is masked from 'package:stats':

loadings

Attaching package: 'MASS'

The following object is masked from 'package:dplyr':

select

Loading required package: Matrix

Attaching package: 'Matrix'

The following objects are masked from 'package:tidyr':

expand, pack, unpack

Loaded glmnet 4.1-1

Loading required package: nlme

Attaching package: 'nlme'

The following object is masked from 'package:dplyr':

collapse

This is mgcv 1.8-35. For overview type 'help("mgcv-package")'.

Registered S3 method overwritten by 'tree':

**method from
print.tree cli**

randomForest 4.6-14

Type rfNews() to see new features/changes/bug fixes.

Attaching package: 'randomForest'

The following object is masked from 'package:dplyr':

combine

The following object is masked from 'package:ggplot2':

margin

Loaded gbm 2.1.8

A data.frame: 500 × 8

logN02	logcars	temp	winds	tempd	windd	hour	day
<dbl>	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>	<int>	<int>
3.71844	7.69120	9.2	4.8	-0.1	74.4	20	600
3.10009	7.69894	6.4	3.5	-0.3	56.0	14	196
3.31419	4.81218	-3.7	0.9	-0.1	281.3	4	513
4.38826	6.95177	-7.2	1.7	1.2	74.0	23	143

4.34640	7.51806	-1.3	2.6	-0.1	65.0	11	115
4.16044	7.67183	2.6	1.6	0.3	224.2	19	527
4.01277	5.52545	-7.9	1.6	0.3	211.9	5	502
2.15176	4.68213	-4.1	3.8	-0.1	63.1	4	453
3.15700	7.15618	-12.7	5.2	-0.1	64.5	12	462
2.37955	4.74493	-1.6	3.0	0.4	58.3	3	554
3.83298	5.81114	-3.1	1.8	0.3	78.0	2	55
4.48187	8.10892	1.0	1.2	1.5	215.0	18	47
4.04830	8.31385	12.2	4.0	-2.8	230.4	17	572
4.00186	5.22036	-1.5	2.4	0.9	82.7	5	556
3.23080	6.40853	-0.9	3.0	0.1	235.0	1	69
4.67189	7.31920	-8.5	0.8	2.9	282.4	20	447
2.73437	6.61740	6.5	4.1	0.0	88.0	24	186
3.49651	7.76938	4.2	7.1	-1.1	9.3	10	550
3.67122	6.46770	-1.1	2.5	1.0	277.0	6	142
3.67377	7.65064	0.8	3.4	-0.4	70.0	12	167
3.15274	7.75061	8.2	4.5	0.2	307.0	14	32
3.42751	5.18178	0.2	0.4	0.2	230.0	5	112
4.32413	7.63964	-2.1	4.3	-0.2	41.4	11	432
3.65584	8.00703	-2.8	6.0	-0.2	62.5	15	453
3.93183	7.75319	-5.4	2.5	1.4	83.0	20	98
2.28238	4.65396	-5.7	4.9	-0.1	58.5	4	409
4.13035	6.67582	6.6	1.6	0.4	104.5	23	579
3.77963	7.79688	5.7	4.8	-1.5	223.0	11	178
3.97968	8.02224	13.8	0.8	-0.1	201.6	17	602
3.75420	7.48829	3.2	4.2	-0.6	285.0	12	121
:	:	:	:	:	:	:	:
2.98062	6.89163	0.1	1.0	0.5	334.0	23	66
3.65066	7.97247	-4.8	3.7	0.0	16.0	18	116
4.25135	7.75018	9.9	1.8	-1.3	266.3	14	529
3.65842	7.42536	4.0	5.8	0.2	184.6	20	402
3.99083	7.63723	-5.4	4.1	-0.2	72.5	17	440


4.44147	6.55108	-12.0	2.8	0.4	82.5	6	444
2.28238	4.98361	1.8	3.5	0.3	121.0	4	128
3.92790	6.57508	5.0	0.8	0.3	262.7	6	403
3.35690	7.09257	4.6	3.4	0.2	194.0	22	109
4.18814	7.83201	-0.2	0.6	-0.3	99.7	7	541
4.20170	7.55329	-0.9	1.1	-0.3	92.0	13	51
2.41591	7.23778	11.9	4.2	-1.3	233.0	15	180
4.33336	8.20084	-9.0	2.4	0.0	84.0	8	100
1.58924	5.81711	1.3	3.8	0.3	33.0	5	174
4.04830	8.26462	7.7	2.6	-0.8	69.1	16	571
2.65324	5.05625	-7.0	3.2	-0.1	67.0	20	85
4.06732	7.73193	2.6	4.2	-0.2	77.9	14	514
3.86283	5.99396	1.6	1.7	1.5	209.1	1	545
3.98155	7.64108	6.7	2.3	-0.4	247.0	10	33
3.81331	7.74587	1.6	2.1	-0.1	146.0	17	167
4.31615	5.14749	-6.5	1.7	0.8	184.0	5	154
3.33932	5.66643	6.7	2.4	-0.1	66.1	5	592
3.96651	7.34923	-0.4	1.9	1.0	77.0	21	72
3.86073	4.45435	-13.8	4.2	0.2	80.0	4	95
2.56495	4.58497	1.8	2.3	0.1	59.0	4	164
4.30946	7.68202	3.5	5.0	-1.0	78.0	11	166
2.94444	6.52942	9.5	6.5	-0.8	210.0	10	35
4.17439	7.75791	5.2	4.6	-0.8	214.0	14	176
2.95491	5.78996	8.4	0.5	-2.6	108.5	7	588
4.03247	8.16223	4.7	5.9	0.4	207.0	17	128

logNO2	logcars	temp	winds
Min. :1.224	Min. :4.127	Min. : -18.6000	Min. :0.300
1st Qu.:3.214	1st Qu.:6.176	1st Qu.: -3.9000	1st Qu.:1.675
Median :3.848	Median :7.425	Median : 1.1000	Median :2.800
Mean :3.698	Mean :6.973	Mean : 0.8474	Mean :3.056
3rd Qu.:4.217	3rd Qu.:7.793	3rd Qu.: 4.9000	3rd Qu.:4.200
Max. :6.395	Max. :8.349	Max. : 21.1000	Max. :9.900


tempd	windd	hour	day
Min. : -5.4000	Min. : 2.0	Min. : 1.00	Min. : 32.0
1st Qu.: -0.2000	1st Qu.: 72.0	1st Qu.: 6.00	1st Qu.:118.8
Median : 0.0000	Median : 97.0	Median :12.50	Median :212.0
Mean : 0.1404	Mean :112.4	Mean :12.28	Mean :210.5

mean	: 0.1494	mean	: 143.4	mean	: 12.38	mean	: 310.5
3rd Qu.:	0.6000	3rd Qu.:	220.0	3rd Qu.:	18.00	3rd Qu.:	513.0
Max.	: 4.3000	Max.	: 359.0	Max.	: 24.00	Max.	: 608.0

logNO2 Versus logcars







logNO2 Versus temp




logNO2 Versus winds




**logNO2 Versus tempd****logNO2 Versus windd**



logNO2 Versus hour



logNO2 Versus day



0	1	2	3	4	5	6					
day											
A data.frame: 500 × 13											
logN02	logcars	temp	winds	tempd	windd	hour	day	wind1	daytime	day7	day:
<dbl>	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>	<int>	<int>	<dbl>	<dbl>	<dbl>	<dbl>
3.71844	7.69120	9.2	4.8	-0.1	74.4	20	600	0	0	5	2
3.10009	7.69894	6.4	3.5	-0.3	56.0	14	196	0	1	0	1
3.31419	4.81218	-3.7	0.9	-0.1	281.3	4	513	1	0	2	1
4.38826	6.95177	-7.2	1.7	1.2	74.0	23	143	0	0	3	1
4.34640	7.51806	-1.3	2.6	-0.1	65.0	11	115	0	1	3	1
4.16044	7.67183	2.6	1.6	0.3	224.2	19	527	1	0	2	1
4.01277	5.52545	-7.9	1.6	0.3	211.9	5	502	1	0	5	1
2.15176	4.68213	-4.1	3.8	-0.1	63.1	4	453	0	0	5	1
3.15700	7.15618	-12.7	5.2	-0.1	64.5	12	462	0	1	0	1
2.37955	4.74493	-1.6	3.0	0.4	58.3	3	554	0	0	1	1
3.83298	5.81114	-3.1	1.8	0.3	78.0	2	55	0	0	6	1
4.48187	8.10892	1.0	1.2	1.5	215.0	18	47	1	1	5	1
4.04830	8.31385	12.2	4.0	-2.8	230.4	17	572	1	1	5	2
4.00186	5.22036	-1.5	2.4	0.9	82.7	5	556	0	0	3	1
3.23080	6.40853	-0.9	3.0	0.1	235.0	1	69	1	0	6	1
4.67189	7.31920	-8.5	0.8	2.9	282.4	20	447	1	0	6	1
2.73437	6.61740	6.5	4.1	0.0	88.0	24	186	0	0	4	1
3.49651	7.76938	4.2	7.1	-1.1	9.3	10	550	0	1	4	1
3.67122	6.46770	-1.1	2.5	1.0	277.0	6	142	1	0	2	1
3.67377	7.65064	0.8	3.4	-0.4	70.0	12	167	0	1	6	1
3.15274	7.75061	8.2	4.5	0.2	307.0	14	32	1	1	4	1
3.42751	5.18178	0.2	0.4	0.2	230.0	5	112	1	0	0	1
4.32413	7.63964	-2.1	4.3	-0.2	41.4	11	432	0	1	5	1
3.65584	8.00703	-2.8	6.0	-0.2	62.5	15	453	0	1	5	1
3.93183	7.75319	-5.4	2.5	1.4	83.0	20	98	0	0	0	1
2.28238	4.65396	-5.7	4.9	-0.1	58.5	4	409	0	0	3	1
4.13035	6.67582	6.6	1.6	0.4	104.5	23	579	0	0	5	2
3.77963	7.79688	5.7	4.8	-1.5	223.0	11	178	1	1	3	1

3.97968	8.02224	13.8	0.8	-0.1	201.6	17	602	1	1	0	2
3.75420	7.48829	3.2	4.2	-0.6	285.0	12	121	1	1	2	1
:	:	:	:	:	:	:	:	:	:	:	:
2.98062	6.89163	0.1	1.0	0.5	334.0	23	66	1	0	3	
3.65066	7.97247	-4.8	3.7	0.0	16.0	18	116	0	1	4	1
4.25135	7.75018	9.9	1.8	-1.3	266.3	14	529	1	1	4	1
3.65842	7.42536	4.0	5.8	0.2	184.6	20	402	1	0	3	
3.99083	7.63723	-5.4	4.1	-0.2	72.5	17	440	0	1	6	
4.44147	6.55108	-12.0	2.8	0.4	82.5	6	444	0	0	3	
2.28238	4.98361	1.8	3.5	0.3	121.0	4	128	0	0	2	1
3.92790	6.57508	5.0	0.8	0.3	262.7	6	403	1	0	4	
3.35690	7.09257	4.6	3.4	0.2	194.0	22	109	1	0	4	1
4.18814	7.83201	-0.2	0.6	-0.3	99.7	7	541	0	0	2	1
4.20170	7.55329	-0.9	1.1	-0.3	92.0	13	51	0	1	2	
2.41591	7.23778	11.9	4.2	-1.3	233.0	15	180	1	1	5	1
4.33336	8.20084	-9.0	2.4	0.0	84.0	8	100	0	1	2	1
1.58924	5.81711	1.3	3.8	0.3	33.0	5	174	0	0	6	1
4.04830	8.26462	7.7	2.6	-0.8	69.1	16	571	0	1	4	2
2.65324	5.05625	-7.0	3.2	-0.1	67.0	20	85	0	0	1	
4.06732	7.73193	2.6	4.2	-0.2	77.9	14	514	0	1	3	1
3.86283	5.99396	1.6	1.7	1.5	209.1	1	545	1	0	6	1
3.98155	7.64108	6.7	2.3	-0.4	247.0	10	33	1	1	5	
3.81331	7.74587	1.6	2.1	-0.1	146.0	17	167	1	1	6	1
4.31615	5.14749	-6.5	1.7	0.8	184.0	5	154	1	0	0	1
3.33932	5.66643	6.7	2.4	-0.1	66.1	5	592	0	0	4	2
3.96651	7.34923	-0.4	1.9	1.0	77.0	21	72	0	0	2	
3.86073	4.45435	-13.8	4.2	0.2	80.0	4	95	0	0	4	
2.56495	4.58497	1.8	2.3	0.1	59.0	4	164	0	0	3	1
4.30946	7.68202	3.5	5.0	-1.0	78.0	11	166	0	1	5	1
2.94444	6.52942	9.5	6.5	-0.8	210.0	10	35	1	1	0	
4.17439	7.75791	5.2	4.6	-0.8	214.0	14	176	1	1	1	1

2.95491	5.78996	8.4	0.5	-2.6	108.5	7	588	0	0	0	2
4.03247	8.16223	4.7	5.9	0.4	207.0	17	128	1	1	2	1

Call:

```
lm(formula = logNO2 ~ logcars, data = z)
```

Residuals:

Min	1Q	Median	3Q	Max
-2.18822	-0.40071	0.06428	0.40362	2.48472

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	1.23310	0.18755	6.575	1.23e-10 ***
logcars	0.35353	0.02657	13.303	< 2e-16 ***


Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 0.6454 on 498 degrees of freedom


Multiple R-squared: 0.2622, Adjusted R-squared: 0.2607

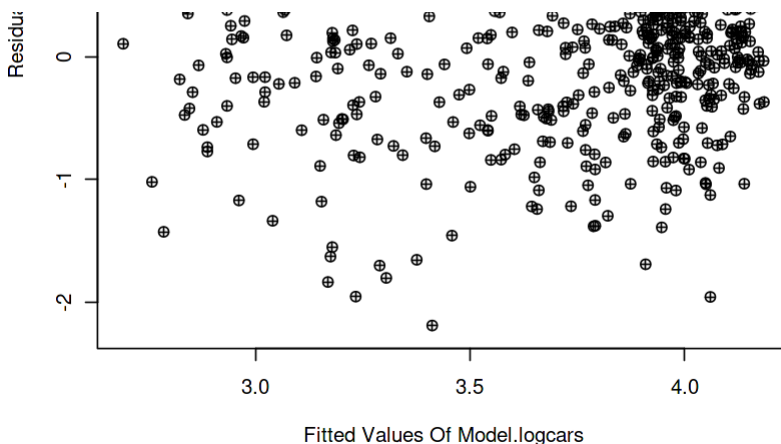
F-statistic: 177 on 1 and 498 DF, p-value: < 2.2e-16

logNO2 Versus day365



Residuals For Model.logcars





Call:

```
lm(formula = logN02 ~ temp, data = z)
```

Residuals:

	Min	1Q	Median	3Q	Max
Residuals	-2.5626	-0.4509	0.1349	0.5037	2.5913

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	3.714761	0.033402	111.215	< 2e-16 ***
temp	-0.019345	0.005082	-3.807	0.000158 ***


Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1


Residual standard error: 0.7407 on 498 degrees of freedom

Multiple R-squared: 0.02828, Adjusted R-squared: 0.02633

F-statistic: 14.49 on 1 and 498 DF, p-value: 0.0001583

Histogram of Model.logcars\$residuals



Residuals For Model.temp

Call:

```
lm(formula = logNO2 ~ winds, data = z)
```

Residuals:

	Min	1Q	Median	3Q	Max
	-2.35783	-0.43975	0.09252	0.49609	2.38463

Coefficients:


	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	4.12113	0.06299	65.43	< 2e-16 ***
winds	-0.13834	0.01780	-7.77	4.51e-14 ***



Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 0.7096 on 498 degrees of freedom

Multiple R-squared: 0.1081, Adjusted R-squared: 0.1063

F-statistic: 60.38 on 1 and 498 DF, p-value: 4.509e-14

Histogram of Model.temp\$residuals



Call:

```
lm(formula = logNO2 ~ tempd, data = z)
```

Residuals:

	Min	1Q	Median	3Q	Max
Residuals	-2.4931	-0.4407	0.1478	0.5050	2.5555

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)	
(Intercept)	3.68003	0.03341	110.140	< 2e-16	***
tempd	0.12273	0.03109	3.947	9.04e-05	***

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1


Residual standard error: 0.7399 on 498 degrees of freedom

Multiple R-squared: 0.03034, Adjusted R-squared: 0.02839


F-statistic: 15.58 on 1 and 498 DF, p-value: 9.04e-05

Histogram of Model.winds\$residuals





Residuals For Model.tempd



Call:

```
lm(formula = logNO2 ~ windd, data = z)
```

Residuals:

Min	1Q	Median	3Q	Max
-2.3641	-0.4642	0.1442	0.5128	2.7519

Coefficients:


	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	3.5574971	0.0646677	55.012	<2e-16 ***
windd	0.0009826	0.0003863	2.544	0.0113 *

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1


Signif. Codes: 0 0.001 0.01 0.05 . 0.1 1

Residual standard error: 0.7465 on 498 degrees of freedom
Multiple R-squared: 0.01282, Adjusted R-squared: 0.01084
F-statistic: 6.47 on 1 and 498 DF, p-value: 0.01127

Histogram of Model.tempd\$residuals



Residuals For Model.windd



Call:

```
lm(formula = logN02 ~ hour, data = z)
```

Residuals:

	Min	1Q	Median	3Q	Max
	-2.2712	-0.4956	0.1000	0.5008	2.6528

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	3.362006	0.067687	49.670	< 2e-16 ***
hour	0.027165	0.004792	5.669	2.44e-08 ***


Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 0.7282 on 498 degrees of freedom


Multiple R-squared: 0.06062, Adjusted R-squared: 0.05873

F-statistic: 32.13 on 1 and 498 DF, p-value: 2.44e-08

Histogram of Model.windd\$residuals



Residuals For Model.hour





Call:

```
lm(formula = logN02 ~ wind1, data = z)
```

Residuals:

	Min	1Q	Median	3Q	Max
Residuals	-2.4121	-0.4581	0.1566	0.5092	2.7592

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	3.63592	0.04545	79.99	<2e-16 ***
wind1	0.13635	0.06716	2.03	0.0429 *


Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 0.7483 on 498 degrees of freedom


Multiple R-squared: 0.008207, Adjusted R-squared: 0.006216


F-statistic: 4.121 on 1 and 498 DF, p-value: 0.04288

Histogram of Model.hour\$residuals



Residuals For Model.wind1





Call:

```
lm(formula = logNO2 ~ daytime, data = z)
```

Residuals:

	Min	1Q	Median	3Q	Max
Residuals	-2.2554	-0.4276	0.1260	0.4577	2.4495

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	3.47917	0.04387	79.298	< 2e-16 ***
daytime	0.46639	0.06400	7.288	1.25e-12 ***


Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 0.7142 on 498 degrees of freedom

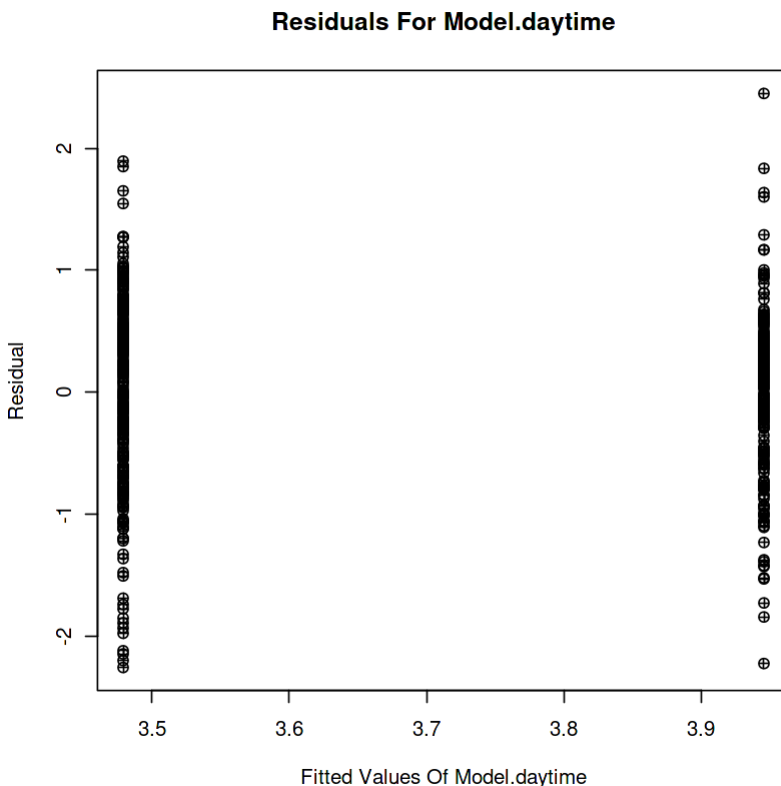
Multiple R-squared: 0.09637, Adjusted R-squared: 0.09455

F-statistic: 53.11 on 1 and 498 DF, p-value: 1.246e-12

Histogram of Model.wind1\$residuals



Model.wind1\$residuals



Call:

lm(formula = logN02 ~ day7, data = z)

Residuals:

	Min	1Q	Median	3Q	Max
Residuals	-2.4009	-0.4728	0.1293	0.5092	2.7212

Coefficients:


	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	3.62469	0.06096	59.463	<2e-16 ***
day7	0.02462	0.01701	1.447	0.148


Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 0.7498 on 498 degrees of freedom

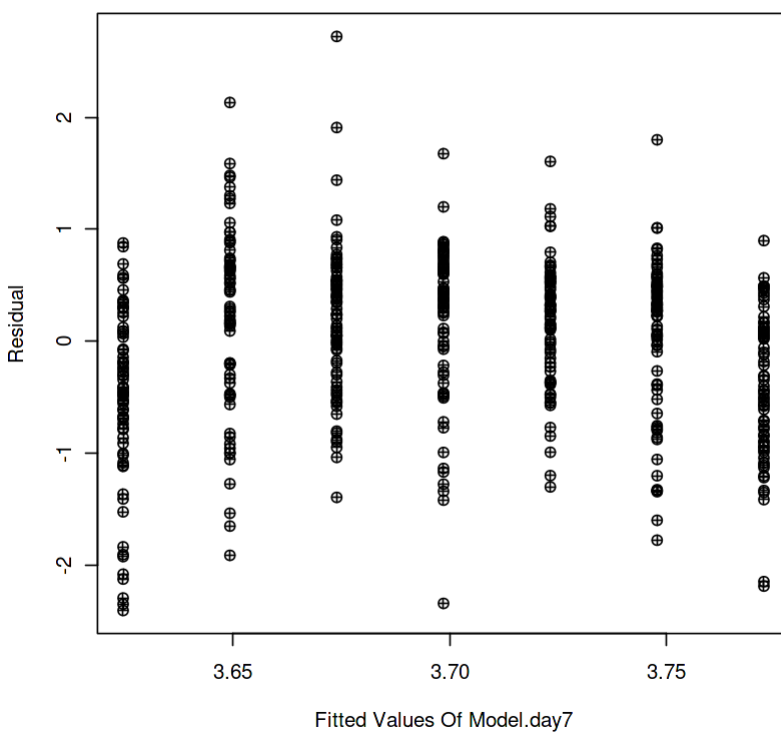
Multiple R-squared: 0.004188, Adjusted R-squared: 0.002189

F-statistic: 2.095 on 1 and 498 DF, p-value: 0.1485

Histogram of Model.daytime\$residuals



Residuals For Model.day7



Call:

```
lm(formula = logNO2 ~ day365, data = z)
```

Residuals:

	Min	1Q	Median	3Q	Max
Residuals	-2.5811	-0.4654	0.1440	0.5353	2.6606

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	3.8546855	0.0793241	48.594	<2e-16 ***
day365	-0.0012146	0.0005589	-2.173	0.0302 *


Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 0.7478 on 498 degrees of freedom


Multiple R-squared: 0.009394, Adjusted R-squared: 0.007405

F-statistic: 4.723 on 1 and 498 DF, p-value: 0.03024

Histogram of Model.day7\$residuals



Residuals For Model.day365



Call:

```
lm(formula = logNO2 ~ daysplit, data = z)
```

Residuals:

Min	1Q	Median	3Q	Max
-2.5008	-0.4554	0.1401	0.4545	2.5333

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	0.00170	0.00000	0.1450	0.8844
daysplit	0.00000	0.00000	0.0000	1.0000

```
(Intercept) 3.86178 0.04222 91.459 < 2e-16 ***
daysplit -0.39857 0.06594 -6.044 2.95e-09 ***
---
```


Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 0.7252 on 498 degrees of freedom


Multiple R-squared: 0.06834, Adjusted R-squared: 0.06647

F-statistic: 36.53 on 1 and 498 DF, p-value: 2.947e-09

Histogram of Model.day365\$residuals



Residuals For Model.daysplit



Call:

```
lm(formula = logNO2 ~ logcars + temp + winds + hour + daytime +
    daysplit, data = z)
```

Residuals:

	Min	1Q	Median	3Q	Max
	-2.11147	-0.31127	-0.00571	0.36681	1.87029

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)	
(Intercept)	1.299074	0.236347	5.496	6.23e-08	***
logcars	0.457403	0.041533	11.013	< 2e-16	***
temp	-0.026084	0.003779	-6.902	1.59e-11	***
winds	-0.146978	0.013763	-10.679	< 2e-16	***
hour	-0.014915	0.004884	-3.054	0.00238	**
daytime	-0.031791	0.074177	-0.429	0.66841	
daysplit	-0.291323	0.048972	-5.949	5.13e-09	***


Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 0.5327 on 493 degrees of freedom


Multiple R-squared: 0.5024, Adjusted R-squared: 0.4963


F-statistic: 82.94 on 6 and 493 DF, p-value: < 2.2e-16

Histogram of Model.daysplit\$residuals



Residuals For Model1





Call:

```
lm(formula = logN02 ~ logcars + temp + winds + hour + daysplit,
   data = z)
```

Residuals:

	Min	1Q	Median	3Q	Max
	-2.09434	-0.30500	-0.00121	0.36370	1.85733

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)	
(Intercept)	1.368052	0.172938	7.911	1.69e-14	***
logcars	0.444086	0.027539	16.126	< 2e-16	***
temp	-0.026068	0.003776	-6.903	1.57e-11	***
winds	-0.147835	0.013605	-10.866	< 2e-16	***
hour	-0.013929	0.004305	-3.236	0.00129	**
daysplit	-0.292934	0.048788	-6.004	3.73e-09	***


Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1



Residual standard error: 0.5323 on 494 degrees of freedom

Multiple R-squared: 0.5022, Adjusted R-squared: 0.4971

F-statistic: 99.66 on 5 and 494 DF, p-value: < 2.2e-16

Histogram of Model1\$residuals



**Residuals For Model2****Histogram of Model2\$residuals**